

Spark SQL Tutorial

Let's make some practice! 😊

How to use Spark?

- **Databricks notebook** community edition - databricks.com
- **Local mode**

install python (<https://www.python.org/downloads/>)

install pip (<https://pip.pypa.io/en/stable/installing/>)

in the **terminal**:

```
pip install findspark
```

then, in the **code**:

```
import findspark
```

```
findspark.init('/path/to/spark')
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.conf import SparkConf
```

```
conf = SparkConf().setAppName(appName).setMaster(master)
```

```
spark = SparkSession.builder.config(conf=conf).getOrCreate()
```

SQL

Dataset

- Distributed collection of data
- Strong typed
- SQL's optimized execution engine
- Only for Scala and Java (but Python already has some features)

DataFrame

- Table in a relational database
- Untyped (Dataset<Row>)
- Many available sources (structured data files, tables, databases, RDDs)
- Scala, Java, Python, and R

Initializing SparkSession

Create a SparkSession (the notebook already has it)

```
from pyspark import SparkSession
spark = SparkSession \
    .builder \
    .appName(appName) \
    .master(master) \
    .getOrCreate()
```

```
sc = spark.sparkContext
```

appName = Name of the application to show on the cluster UI

master = Spark URL ('spark://ip-address:7077') or 'local'

Creating DataFrames

```
df = spark.read.csv('files/people.csv')
```

`df.show()` outputs the content of the dataframe

`df.printSchema()` outputs the schema of the dataframe

Untyped Dataset Operations (aka DataFrame Operations)

// Select only the "name" column

```
df.select("name").show()
```

// Select only the "name" and "address" column and add 1 to "age"

```
df.select(col("name"), col("address"), col("age").plus(1)).show()
```

In order to access nested element do

root.child e.g. "address.city"

Untyped Dataset Operations (aka DataFrame Operations)

```
// Select people older than 21
```

```
df.filter(col("age").gt(21)).show()
```

```
// Count people by age
```

```
df.groupBy("age").count().show()
```

Running SQL Queries Programmatically

```
// Register the DataFrame as a SQL temporary view
```

```
df.createOrReplaceTempView("people")
```

```
sqlDF = spark.sql("SELECT * FROM people")
```

```
sqlDF.show()
```


Running SQL Queries Programmatically

Temporary views are session-scoped -> they disappear after termination
In order to keep it alive, create a global temporary view

```
spark = spark.newSession() // table people is no longer present
```

```
df.createGlobalTempView("people")
```

```
spark.newSession().sql("SELECT * FROM global_temp.people").show();
```

Creating Datasets

```
public class Person implements Serializable {  
    private String name;  
    private int age;  
    private Address address;  
  
    /** Getters and Setters */  
}
```

```
public class Address implements Serializable {  
    private String city;  
    private String state;  
  
    /** Getters and Setters */  
}
```

Creating Datasets

```
Person person = new Person();  
person.setName("Andy");  
person.setAge(32);  
Address address = new Address();  
address.setCity("Rome");  
address.setState("Italy");  
person.setAddress(address);
```

```
Encoder<Person> personEncoder = Encoders.bean(Person.class);  
Dataset<Person> dataset =  
    spark.createDataset(Collections.singletonList(person), personEncoder);  
  
dataset.show();
```

Creating Datasets

```
Encoder<Integer> integerEncoder = Encoders.INT();
```

```
Dataset<Integer> primitiveDS =  
    spark.createDataset(Arrays.asList(1, 2, 3), integerEncoder);
```

```
Dataset<Integer> transformedDS = primitiveDS.map(  
    (MapFunction<Integer, Integer>) value -> value + 1,  
    integerEncoder);
```

```
transformedDS.collect();
```

Creating Datasets

```
Encoder<Person> personEncoder = Encoders.bean(Person.class);
```

```
Dataset<Person> personDataset =  
    spark.read().json("files/people.json").as(personEncoder);
```

```
personDataset.show();
```

Interoperating with RRDs

Inferring the Schema Using Reflection

Concise syntax, schema already known

```
lines = sc.textFile("examples/src/main/resources/people.txt")
parts = lines.map(lambda l: l.split(","))
people = parts.map(lambda p: Row(name=p[0], age=int(p[1])))
```

Interoperating with RRDs

Inferring the Schema Using Reflection

```
schemaPeople = spark.createDataFrame(people)
schemaPeople.createOrReplaceTempView("people")
```

```
teenagers = spark.sql("""
    SELECT name FROM people
    WHERE age >= 13 AND age <= 19""")
```

```
teenNames = teenagers.rdd.map(lambda p: "Name: " + p.name).collect()
```

```
for name in teenNames:
    print(name)
```

Interoperating with RRDs

Programmatically Specifying the Schema

Verbose syntax, schema not known

```
lines = sc.textFile("examples/src/main/resources/people.txt")\
    .map(lambda l: l.split(","))\
    .map(lambda p: (p[0], p[1].strip()))
```

```
schemaStr = "name age city state";
```

```
fields = [StructField(field, StringType(), True) for field in schemaStr.split()]
```

```
schema = StructType(fields)
```


Interoperating with RRDs

Programmatically Specifying the Schema

```
schemaPeople = spark.createDataFrame(people, schema)
```

```
schemaPeople.createOrReplaceTempView("people")
```

```
results = spark.sql("SELECT name FROM people")
```

```
results.show()
```

Data Sources

Generic Load/Save Functions

Default: Parquet file

```
usersDF = spark.read.load("files/users.parquet")
```

```
usersDF.select("name", "favorite_color").write.save("files/results.parquet")
```

```
sqlDF = spark.sql("SELECT * FROM parquet.`files/users.parquet`")
```

Exercises [both RDD and SparkSQL]

- **Use** github.com/forons/BigDataExamples/blob/master/files/tweets_cleaned.csv
- **Count tweets per user**
- **Split created_at field and count the number of tweets per hour**
- **Count the tweets that contain a word that you choose**
- **Sort the users based on the number of tweets**

Exercises [both RDD and SparkSQL]

- Find all the tweets by user
- Find how many tweets each user has
- Find all the persons mentioned on tweets
- Count how many times each person is mentioned
- Find the 10 most mentioned persons
- Find all the hashtags mentioned on a tweet
- Count how many times each hashtag is mentioned
- Find the 10 most popular Hashtags

Contacts

For any problem, send a mail to

daniele.foroni@unitn.it